

ARMv8 Assembler Language

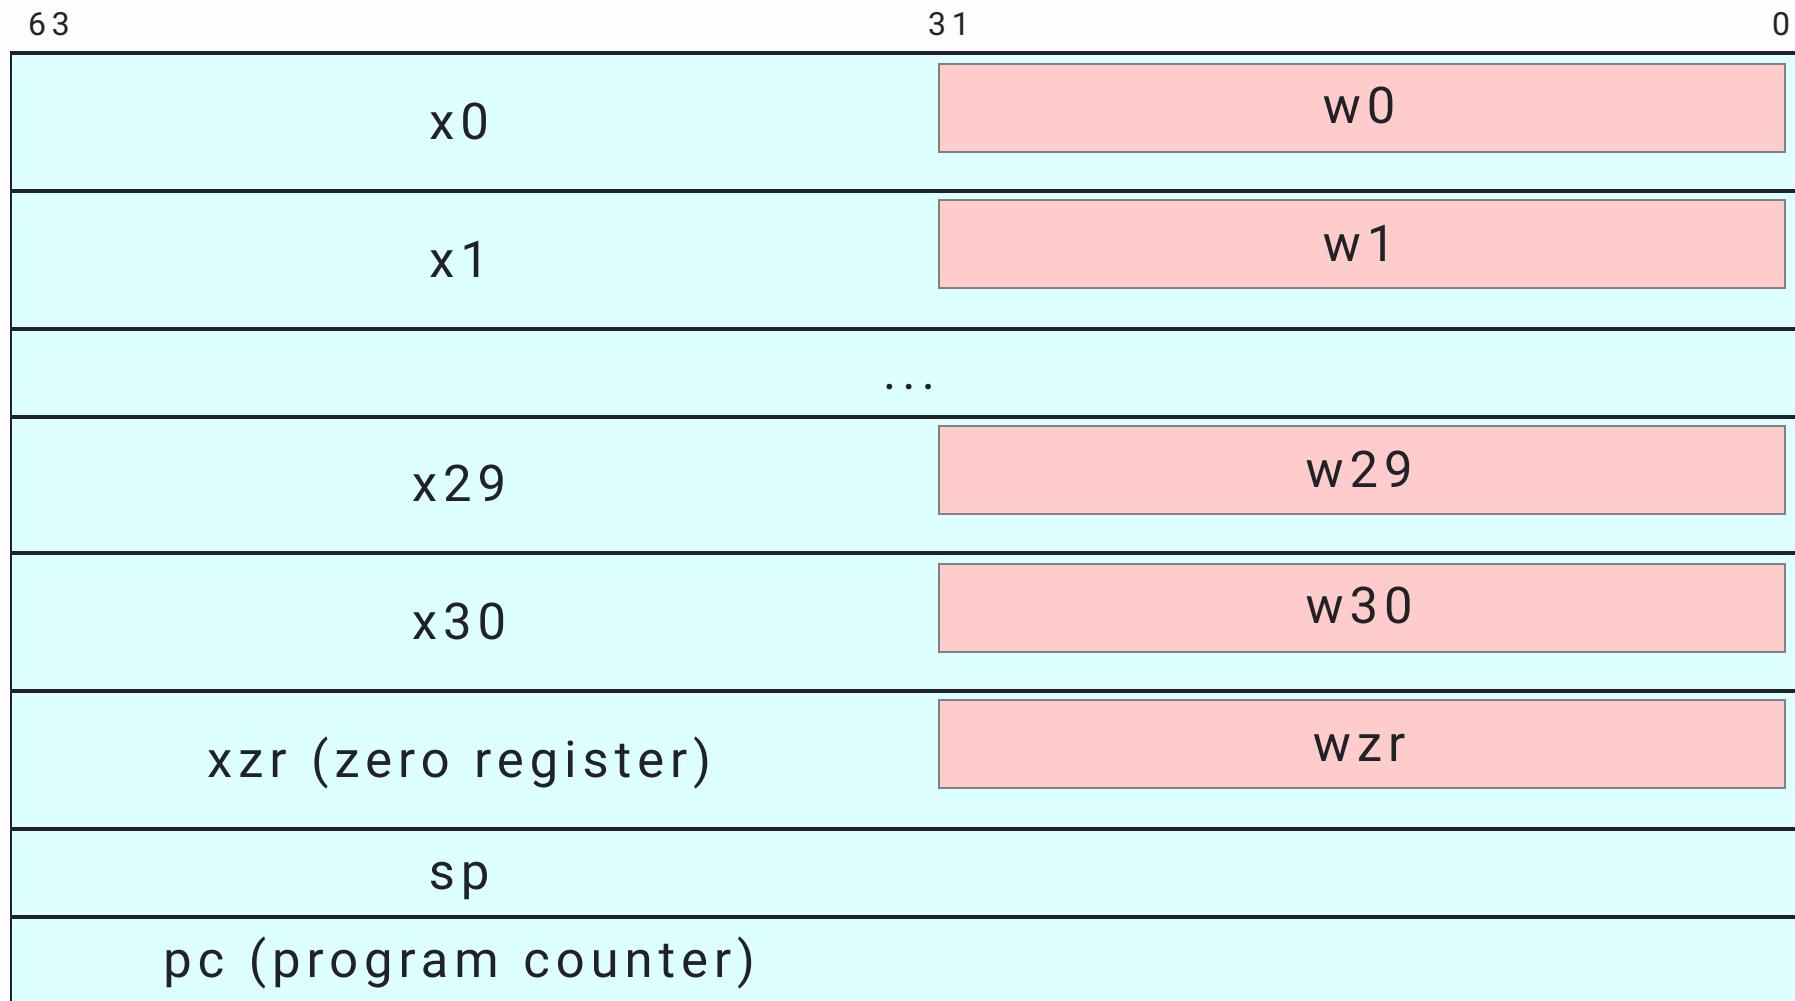
Tutorial

Mike Rogers

Introduction

- RISC pattern is Load/Store:
 - Load data from RAM to registers
 - Manipulate data in registers
 - Store data from registers to RAM
- On AARCH64 (also known as ARMv8) is RISC
 - Math and logic instructions can only access registers

Registers in the AARCH64



 N Z C V D A I F FIQ SS IL EL nRW SP Q GE IT J T E M

General-Purpose Registers

- Registers X0 through X30
- these are 64-bit registers
- Used to hold data for instructions
- Used to pass parameters to and from functions
- Used to hold the return address for function calls (X0)
- Some have special purposes defined in hardware (e.g. X30) or defined by software convention (e.g. X29)
- Register W0 through W30 are the 32 bit versions, and take up the 32 least significant bits of the 64 bit registers

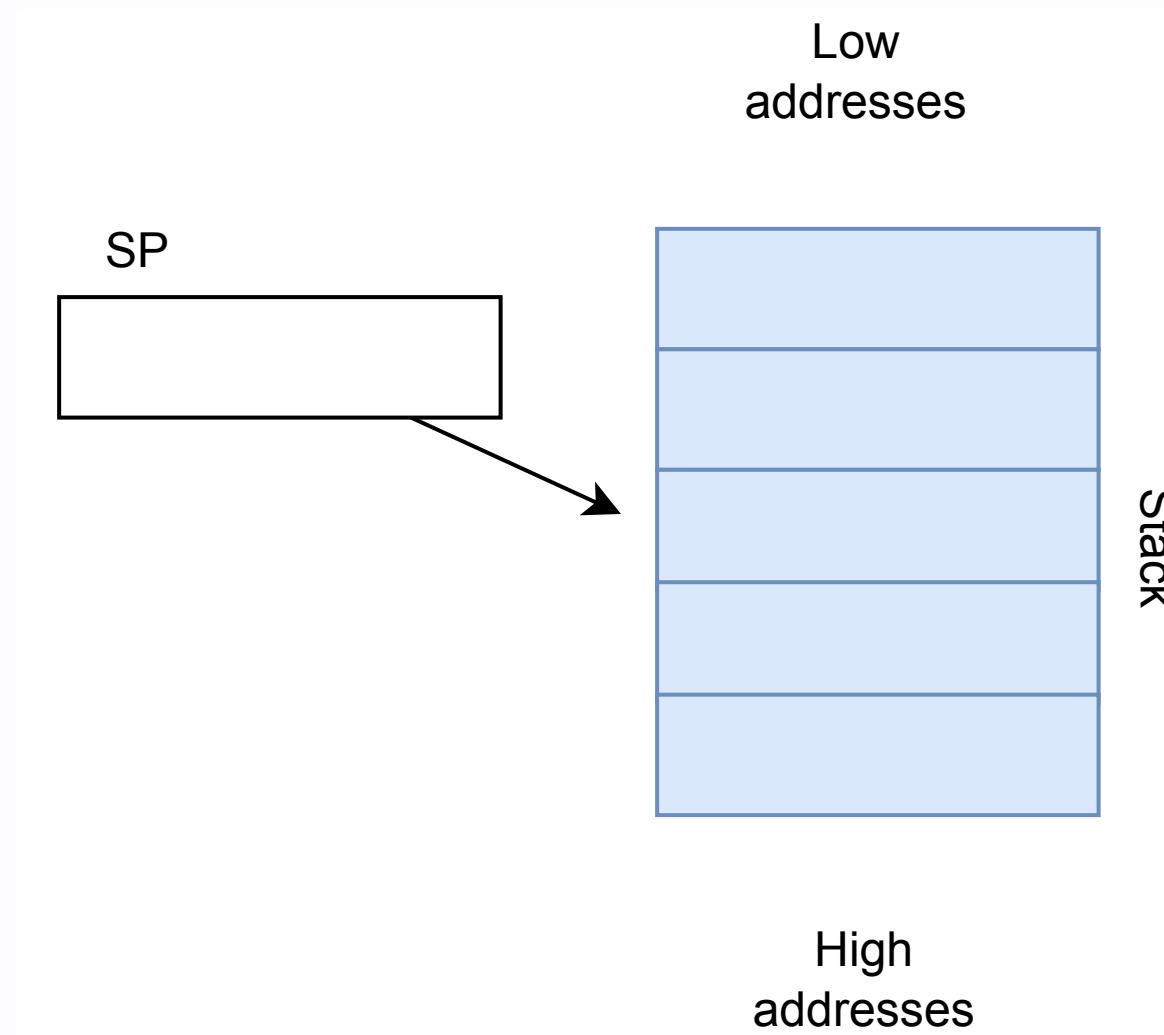
Zero Register

- The Zero register is XZR
- Reading it gives all zeros
- Any data written to it is thrown away

The Stack Pointer Register

- The stack pointer register **SP** is a special-purpose register
- It contains the *Stack Pointer*:
 - which is the address of top (low address) of current function's stack frame
 - Stack grow toward low addresses

The Stack Pointer Register (2)



PSTATE

- The processor state, or PSTATE, is not held in a single register
- And is not accessible directly
- It is a set of bits that represent...
 - condition flags set from instruction execution:
 - n (Negative), z (Zero), c (Carry), v (oVerflow)
 - Programs often use the compare (cmp) instruction to set flags
 - and then use conditional branch instructions to execute the appropriate code
 - beq, bne, blo, bhi, ble, bge, ...

Instruction Format

- The general format of AARCH64 instructions is:

```
instr_name dest_reg, src_reg1, src_reg2
```

```
instr_name dest_reg, src_reg1, immediate
```

- `dest_reg`, `src_reg1`, and `src_reg2` are all 64-bit `x` registers for 64-bit operations
- `dest_reg`, `src_reg1`, and `src_reg2` are all 32-bit `w` registers for 32-bit operations
- Note that you cannot "mix" register sizes

64 bit Math

- C example:

```
static long long i = 20;
static long long j = 30;
static long long k = 40;
int main() {
    k = (i + j) * 2;
    //...
}
```

- How does i,j, and k get initialized? via compiler directives (later)

- AARCH64 Assembly

```
// move address of i into x0 and
// move value of i into x0
adr    x0, i
ldr    x0, [x0]
// move address of j into x1 and
// move value of j into x1
adr    x1, j
ldr    x1, [x1]
// add i and j, put result into x2 and
// left shift by 1 = mult by 2
add    x2, x0, x1
lsl    x2, x2, 1
// move address of k into x3 and
// store result into k
adr    x3, k
str    x2, [x3]
```

Another Math Example

- C example

```
static long long x;
static long long y;
static long long z;

int main() {
    z = x - y;
    z = x * y;
    z = x / y;
    z = x & y;
    z = x | y;
    z = x ^ y;
    z = x >> y;

    return z;
}
```

Another Math Example (2)

```

adr    x0, x      // load x and y
ldr    x1, [x0]
adr    x0, y
ldr    x0, [x0]
sub   x1, x1, x0 // z = x - y;
adr    x0, z
str   x1, [x0]
adr    x0, x      // load x and y
ldr    x1, [x0]
adr    x0, y
ldr    x0, [x0]
mul   x1, x1, x0 // z = x * y;
adr    x0, z
str   x1, [x0]
adr    x0, x      // load x and y
ldr    x1, [x0]
adr    x0, y
ldr    x0, [x0]
sdiv  x1, x1, x0 // z = x / y
adr    x0, z
str   x1, [x0]
adr    x0, x      // load x and y
ldr    x1, [x0]
adr    x0, y
ldr    x0, [x0]
and   x1, x1, x0 // z = x & y
adr    x0, z
str   x1, [x0]
  
```

```

adr    x0, x      // load x and y
ldr    x1, [x0]
adr    x0, y
ldr    x0, [x0]
orr   x1, x1, x0 // z = x | y;
adr    x0, z
str   x1, [x0]
adr    x0, x      // load x and y
ldr    x1, [x0]
adr    x0, y
ldr    x0, [x0]
eor   x1, x1, x0 // z = x ^ y
adr    x0, z
str   x1, [x0]
adr    x0, x      // // load x and y
ldr    x0, [x0]
adr    x1, y
ldr    x1, [x1]
asr   x1, x0, x1 // z = x >> y
adr    x0, z
str   x1, [x0]
adr    x0, z      // return z
ldr    x0, [x0]
ret
  
```

Instructions that save typing

```
mov x3, x1 // same as: orr x3, xzr, x1
neg x3, x1 // same as: sub x3, xzr, x1
```

- remember, xzr is the zero register

Signed and Unsigned Math

- C code:

```
unsigned long long x = 20;  
long long y = 20;  
x++;  
y--;
```

- AARCH64 assembly

| | |
|-----|-----------|
| adr | x0, x |
| ldr | x1, [x0] |
| add | x1, x1, 1 |
| str | x1, [x0] |

| | |
|-----|-----------|
| adr | x0, y |
| ldr | x1, [x0] |
| sub | x1, x1, 1 |
| str | x1, [x0] |

- The same instructions
 - However, sets different condition flags in PSTATE
- Exception is division: `sdiv` vs `udiv` instructions

32-bit Arithmetic

- C code

```
static int length = 20;
static int width = 40;
static int perim;
int main() {
    perim = (length + width) * 2;
    return perim;
}
```
 - AARCH64 code

```
adr      x0, length
adr      x1, width
adr      x2, perim
ldr      w0, [x0]
ldr      w1, [x1]
add      w0, w0, w1
lsl      w0, w0, 1 // mult by 2
str      w0, [x2]
```
- Note that addresses are still 64 bit, so you must use the Xn registers (in brackets)

8- and 16-bit Arithmetic

- C Code

```
static short int x = 20;
static short int y = 30;
static short int z;
int main() {
    z = x + y;
    return z;
}
```

- AARCH64 code

| | |
|-------|------------|
| adr | x0, x |
| adr | x1, y |
| adr | x2, z |
| ldrsh | w0, [x0] |
| ldrsh | w1, [x1] |
| add | w0, w0, w1 |
| strh | w0, [x2] |

- strh is *store register halfword*
- Use special *load* and *store* instructions for transfer of shorter data types from and to memory
- more later

Loading and storing

```
ldr dest, [src]  
str src, [dest]
```

- dest and src must be registers
- Registers in [brackets] contain memory addresses (i.e. pointer)

Loading and storing (2)

- C code

```
static int length = 1;
static int width = 2;
static int perim = 0;
int main() {
    perim = (length + width) * 2;
    return 0;
}
```

- AARCH64 code

```
.section data
length:    .int 1
width:     .int 2
perim:     .int 0

.section      .text
.global        main
main:
    adr      x0, length
    ldr      w1, [x0]
    adr      x0, width
    ldr      w2, [x0]
    add      w1, w1, w2
    lsl      w1, w1, 1
    adr      x0, perim
    str      w1, [x0]
    mov      w0, 0
    ret
```

Sections

- Declaring sections in GAS assembly
 - `.data` is a read-write data section of initialized data
 - `.rodata` is a read-only section of initialized data
 - `.bss` is a read-write section of data initialized to zero (zeroed when loaded)
 - `.text` is execute-only, program code
- More later on the stack

Memory Labels

- Humans do not like remembering addresses for memory
- But, we like names - they are easy to remember
- *Labels* are names for locations in memory
- Example

```
.section .data
my_int:      .word 1234 // .word and .int are the same size for AARCH64
my_int2:     .int  1234
```

Global Symbols

- Declare `main` to be a globally-visible label
- Globally visible labels can be used outside of the file in which they are defined
- So, global variables and functions can be *exported*

```
.section .data
length:    .int 1
width:     .int 2
perim:      .int 0
.section .text
.global main
main:
    adr    x0, length
    ldr    w1, [x0]
    adr    x0, width
    ldr    w2, [x0]
    add    w1, w1, w2
    lsl    w1, w1, 1
    adr    x0, perim
    str    w1, [x0]
    mov    w0, 0
    ret
```

Addressing

```
.section .data
length: .int 1
width: .int 2
perim: .int 0
.section .text
.global main
main:
    adr x0, length
    ldr w1, [x0]
    adr x0, width
    ldr w2, [x0]
    add w1, w1, w2
    lsl w1, w1, 1
    adr x0, perim
    str w1, [x0]
    mov w0, 0
    ret
```

- `adr` is an instruction that loads an address of a label into a register

Loading and Storing with Addresses

```
.section .data
length: .int 1
width: .int 2
perim: .int 0
.section .text
.global main
main:
    adr x0, length
    ldr w1, [x0]
    adr x0, width
    ldr w2, [x0]
    add w1, w1, w2
    lsl w1, w1, 1
    adr x0, perim
    str w1, [x0]
    mov w0, 0
    ret
```

- Once the address is loaded into `x0`, use the address (or *pointer*) in `x0` to load the data from (or store it to) memory

Trace

| | | | |
|----|----------------|--------|---|
| x0 | addr of length | length | 1 |
| w1 | ? | width | 2 |
| w2 | ? | perim | 0 |

```
.section .data
length:    .int 1
width:     .int 2
perim:     .int 0
.section .text
.global main
main:
    →adr    x0, length
    ldr    w1, [x0]
    adr    x0, width
    ldr    w2, [x0]
    add    w1, w1, w2
    lsl    w1, w1, 1
    adr    x0, perim
    str    w1, [x0]
    mov    w0, 0
    ret
```

Trace (2)

| | | | |
|----|----------------|--------|---|
| x0 | addr of length | length | 1 |
| w1 | 1 | width | 2 |
| w2 | ? | perim | 0 |

```
.section .data
length:    .int 1
width:     .int 2
perim:     .int 0
.section .text
.global main
main:
    adr    x0, length
    → ldr    w1, [x0]
    adr    x0, width
    ldr    w2, [x0]
    add    w1, w1, w2
    lsl    w1, w1, 1
    adr    x0, perim
    str    w1, [x0]
    mov    w0, 0
    ret
```

Trace (3)

| | | | |
|----|---------------|--------|---|
| x0 | addr of width | length | 1 |
| w1 | 1 | width | 2 |
| w2 | ? | perim | 0 |

```
.section .data
length:    .int 1
width:     .int 2
perim:     .int 0
.section .text
.global main
main:
    adr    x0, length
    ldr    w1, [x0]
→   adr    x0, width
    ldr    w2, [x0]
    add    w1, w1, w2
    lsl    w1, w1, 1
    adr    x0, perim
    str    w1, [x0]
    mov    w0, 0
    ret
```

Trace (4)

| | | | |
|----|---------------|--------|---|
| x0 | addr of width | length | 1 |
| w1 | 1 | width | 2 |
| w2 | 2 | perim | 0 |

```
.section .data
length:    .int 1
width:     .int 2
perim:     .int 0
.section .text
.global main
main:
    adr    x0, length
    ldr    w1, [x0]
    adr    x0, width
    →ldr   w2, [x0]
    add    w1, w1, w2
    lsl    w1, w1, 1
    adr    x0, perim
    str    w1, [x0]
    mov    w0, 0
    ret
```

Trace (5)

| | | | |
|----|---------------|--------|---|
| x0 | addr of width | length | 1 |
| w1 | 3 | width | 2 |
| w2 | 2 | perim | 0 |

```
.section .data
length:    .int 1
width:     .int 2
perim:     .int 0
.section .text
.global main
main:
    adr    x0, length
    ldr    w1, [x0]
    adr    x0, width
    ldr    w2, [x0]
→add   w1, w1, w2
    lsl   w1, w1, 1
    adr    x0, perim
    str   w1, [x0]
    mov   w0, 0
    ret
```

Trace (6)

| | | | |
|----|---------------|--------|---|
| x0 | addr of width | length | 1 |
| w1 | 6 | width | 2 |
| w2 | 2 | perim | 0 |

```
.section .data
length:    .int 1
width:     .int 2
perim:     .int 0
.section .text
.global main
main:
    adr    x0, length
    ldr    w1, [x0]
    adr    x0, width
    ldr    w2, [x0]
    add    w1, w1, w2
    →lsl   w1, w1, 1
    adr    x0, perim
    str    w1, [x0]
    mov    w0, 0
    ret
```

Trace (7)

| | | | |
|----|---------------|--------|---|
| x0 | addr of perim | length | 1 |
| w1 | 6 | width | 2 |
| w2 | 2 | perim | 0 |

```
.section .data
length:    .int 1
width:     .int 2
perim:     .int 0
.section .text
.global main
main:
    adr    x0, length
    ldr    w1, [x0]
    adr    x0, width
    ldr    w2, [x0]
    add    w1, w1, w2
    lsl    w1, w1, 1
    →adr  x0, perim
    str    w1, [x0]
    mov    w0, 0
    ret
```

Trace (8)

| | | | |
|----|---------------|--------|---|
| x0 | addr of perim | length | 1 |
| w1 | 6 | width | 2 |
| w2 | 2 | perim | 6 |

```
.section .data
length:    .int 1
width:     .int 2
perim:     .int 0
.section .text
.global main
main:
    adr    x0, length
    ldr    w1, [x0]
    adr    x0, width
    ldr    w2, [x0]
    add    w1, w1, w2
    lsl    w1, w1, 1
    adr    x0, perim
    →str   w1, [x0]
    mov    w0, 0
    ret
```

Cleanup

```
.section .data
length:    .int 1
width:     .int 2
perim:     .int 0
.section .text
.global main
main:
    adr    x0, length
    ldr    w1, [x0]
    adr    x0, width
    ldr    w2, [x0]
    add    w1, w1, w2
    lsl    w1, w1, 1
    adr    x0, perim
→   str    w1, [x0]
→   mov    w0, 0
    ret
```

- Note that the return value is put in W0
- and we return with `ret`

More addressing modes

- A64 has more addressing modes
- These modes allow you to flexibly manipulate addresses, and are particularly useful when accessing arrays
 - in loops
- However, as this is a beginning tutorial, these addressing modes will not be covered in these slides

More on defining data

- C code

```
static char c = 'Q'; // 8 bits
static short s = 7; // 16 bits
static int i = 777; // 32 bits
static long l = 888; // also 32 bits
static long long ll = 999; // 64 bits
```

- ARM code

```
.section .data
c:    .byte  'Q' // 8 bits
s:    .short 7 // 16 bits
i:    .word 777 // 32 bits
          // or .int or .long
l:    .int 888 // like this
ll:   .quad 999 // 64 bits
str:  .asciz 'Hello World' // a string
```

Assembler Global Variables

```
.section .data
.global c // c and s can be used from code in
.global s // other files, but not i, l, or ll
c:    .byte  'Q'   // 8 bits
s:    .short 7    // 16 bits
i:    .word  777   // 32 bits
          // or .int or .long
l:    .int   888   // like this
ll:   .quad  999   // 64 bits
str:  .asciz 'Hello World' // a string
```

- Use `.global` to export global variables define in assembly for use from other files (e.g. your C program)

What IF

- C example

```
static int i = 9;
int main () {
    if (i == 9) {
        i = 0;
    }
    return i;
}
```

- How do you turn this into assembly?

What IF (2)

- C example

```
static int i = 9;
int main () {
    if (i != 9) goto skipit;
    i = 0;
skipit:
    return i;
}
```

- First, make it easier to deal with by only having a `goto` in the body

What IF (3)

- Now, turn it into assembly

```
.section .data
i: .int 9
.section .text
.global main
main:
    adr      x0, i
    ldr      w1, [x0]
    cmp      w1, 9          // subs wzr, w1, 9
    bne      skipit
    mov      w1, 0
    str      w1, [x0]
skipit:
    mov      w0, w1
ret
```

IF...ELSE Example

```

int i = 10;
int j = 20;
int smaller;
int main() {
    if (i < j)
        smaller = i;
    else
        smaller = j;
}
  
```

```

// ...
    adr    x0, i
    ldr    w1, [x0]
    adr    x0, j
    ldr    w2, [x0]
    cmp    w1, w2 // if !(i < j)
    bge    else //      goto else
    adr    x0, smaller
    str    w1, [x0]
    b      endif
else: // i is _not_ less than j
    adr    x0, smaller // j is smaller
    str    w2, [x0]
endif:
  
```

WHILE Example

```
static int n = 20;
static int fact = 0;
int main() {
    fact = 1;
    while (n > 1) {
        fact *= n;
        n--;
    }
    return fact;
}
```

```
// ...
adr x0, n
ldr w1, [x0]
mov w2, 1
while:
    cmp w1, 1 // !(n > 1)
    ble endwhile
    mul w2, w2, w1
    sub w1, w1, #1 // the # is optional
                    // for immediates
    b while
endwhile:
// ...
```

Branch Modifiers

Table 3.2: AArch64 condition modifiers.

| Condition code | Meaning | Condition flags | Binary encoding |
|----------------|------------------------------------|---------------------------|-----------------|
| EQ | Equal | $Z = 1$ | 0000 |
| NE | Not Equal | $Z = 0$ | 0001 |
| HI | Unsigned Higher | $(C = 1) \wedge (Z = 0)$ | 1000 |
| HS | Unsigned Higher or Same | $C = 1$ | 0010 |
| LS | Unsigned Lower or Same | $(C = 0) \vee (Z = 1)$ | 1001 |
| LO | Unsigned Lower | $C = 0$ | 0011 |
| GT | Signed Greater Than | $(Z = 0) \wedge (N = V)$ | 1100 |
| GE | Signed Greater Than or Equal | $N = V$ | 1010 |
| LE | Signed Less Than or Equal | $(Z = 1) \vee (N \neq V)$ | 1101 |
| LT | Signed Less Than | $N \neq V$ | 1011 |
| CS | Unsigned Overflow (Carry Set) | $C = 1$ | 0010 |
| CC | No Unsigned Overflow (Carry Clear) | $C = 0$ | 0011 |
| VS | Signed Overflow | $V = 1$ | 0110 |
| VC | No Signed Overflow | $V = 0$ | 0111 |
| MI | Minus, Negative | $N = 1$ | 0100 |
| PL | Plus, Positive or Zero | $N = 0$ | 0101 |
| AL | Always Executed | Any | 1110 |
| NV | Never Executed | Any | 1111 |

Calling functions

- Use `bl` (branch and link)
 - stores return point in X30
- `ret` (return) instruction returns to address in X30

```
int main() {
    f();
}
void f() {
    // do something
}
```

```
// ...
main:
    bl      f
            mov      w0, 0
            ret
f:
    // do something
    ret
```

Nested Function Calls

- But what about this?

```
int main() {  
    f();  
    return 0;  
}  
  
void f() {  
    g();  
}  
  
void g() {  
    // do something  
}
```

```
main:  
    bl      f  
    mov     x0, 0  
    ret  
  
f:  
    bl      g  
    ret  
  
g:  
    // do something  
    ret
```

- This will not work. Why?

Use the stack

- Push X30 on stack when entering a function
- Pop X30 from stack before returning from a function

main:

```
// Save X30
sub sp, sp, 16
str x30, [sp]
bl f
// Restore X30
ldr x30, [sp]
add sp, sp, 16
ret
```

...

f:

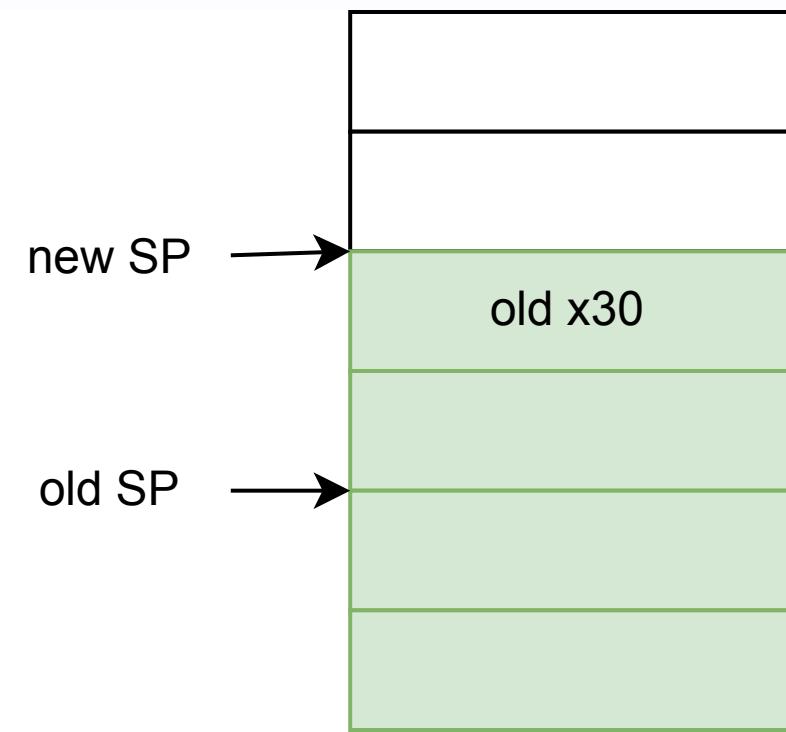
```
// Save X30
sub sp, sp, 16
str x30, [sp]
bl g
// Restore X30
ldr x30, [sp]
add sp, sp, 16
ret
```

The Stack

- SP (*stack pointer*) register points to top of stack
- It can be used in ldr and str instructions
- It can be used in arithmetic instructions
- However, the address in SP *must* be multiple of 16
- That's why, even though in the previous example we only needed 8 bytes to store x30...
 - we still needed to grow the stack by 16 bytes

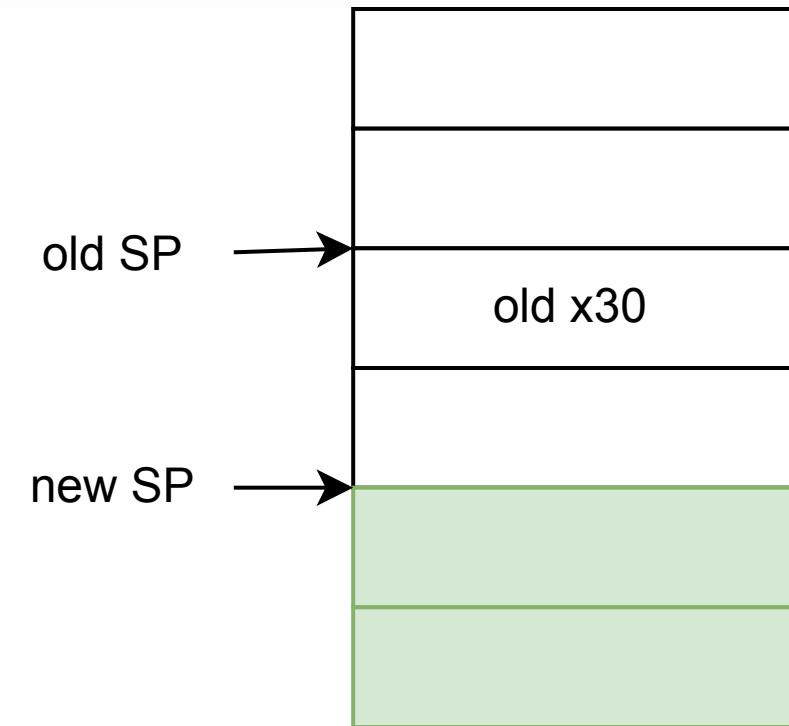
The Stack (2)

```
sub sp, sp, 16  
str x30, [sp]
```



The Stack (2)

```
add sp, sp, 16
```



Passing Arguments

- Pass first 8 (integer or address) arguments in registers
 - X0..X7 and/or W0..W7
- More than 8 arguments
 - Pass them on the stack
- if arguments are structures (pass by value)
 - Pass them on the stack

Passing Arguments Example

```
// int do_something(int x, int y);
do_something:
    sub    sp, sp, 32      // Why 32? Because we will store
                          // three values on the stack
                          // x30 (link), x, and y.
                          // remember, SP must be a multiple
                          // of 16
    str    x30, [sp]        // Save x30
    str    x0, [sp, 16]      // Save x
    str    x1, [sp, 8]       // Save y

    // now we can use x0 and x1 as scratch registers
    // or to pass parameters to functions we call.
    // We load x and y when we need them.
    // x30 is also saved, which make calling other functions
    // safe.

    //...

    ldr    x30, [sp]
    // Restore x30
    add    sp, sp, 32
    ret
```

What Registers Do I Save?

Next, the PCS defines which registers can be corrupted, and which registers cannot be corrupted. If a register can be corrupted, then the called function can overwrite without needing to restore, as this table of PCS register rules shows:

| X0-X7 | X8-X15 | X16-X23 | X24-X30 |
|--|--------------------------------|----------------------------------|----------------------------------|
| Parameter and Result Registers (x0-x7) | xR (x8) | IP0 (x16) | Callee-saved Registers (x24-x28) |
| - | Corruptible Registers (x9-x15) | IP1 (x17) | FP (x29) |
| - | - | PR (x18) | LR (x30) |
| - | - | Callee-saved Registers (x19-X23) | - |

For example, the function `foo()` can use registers `x0` to `x15` without needing to preserve their values. However, if `foo()` wants to use `x19` to `x28` it must save them to stack first, and then restore from the stack before returning.

Some registers have special significance in the PCS:

- `xR` - This is an indirect result register. If `foo()` returned a struct, then the memory for struct would be allocated by the caller, `main()` in the earlier example. `xR` is a pointer to the memory allocated by the caller for returning the struct.
- `IP0` and `IP1` - These registers are intra-procedure-call corruptible registers. These registers can be corrupted between the time that the function is called and the time that it arrives at the first instruction in the function. These registers are used by linkers to insert veneers between the caller and callee. Veneers are small pieces of code. The most common example is for branch range extension. The branch instruction in A64 has a limited range. If the target is beyond that range, then the linker needs to generate a veneer to extend the range of the branch.
- `FP` - Frame pointer.
- `LR` - `x30` is the link register (`LR`) for function calls.